

Random Gradient-Free Minimization of Convex Functions

A 20-minute presentation on zeroth-order convex optimization

Yurii Nesterov and Vladimir Spokoiny

ECE 490 Final Project

Presenter	Main responsibility
Zhihao Jin	smoothing, oracle theory, acceleration
Yuanxiang Gong	nonsmooth and stochastic optimization
Girish Ravi	smooth baseline, nonconvex extension, experiments, closing


Motivation: Why Zeroth-Order Optimization?

Paper Sec. 1.1 — Owner: Leader

$$\min_{x \in \mathbb{R}^n} f(x)$$

- Classical methods need $\nabla f(x)$.
- Zeroth-order methods only query $f(x)$.
- Useful for black-box simulations, legacy code, memory-heavy AD, or nonsmooth models.

Question: what is the price of not using gradients?




Gaussian Smoothing: The Bridge to Gradients

Paper Sec. 2 — Owner: Leader

$$f_\mu(x) = \mathbb{E}_u[f(x + \mu u)], \quad u \sim \mathcal{N}(0, I).$$

$$\nabla f_\mu(x) = \mathbb{E}_u \left[\frac{f(x + \mu u) - f(x)}{\mu} u \right].$$

- Finite differences estimate ∇f_μ , not always ∇f .
- Nonsmooth bias: $|f_\mu(x) - f(x)| \leq \mu L_0 \sqrt{n}$.
- Smooth bias: $|f_\mu(x) - f(x)| \leq \frac{\mu^2}{2} L_1 n$.



Random Oracles and the Dimension Factor

Paper Sec. 3 — Owner: Leader


Euclidean notation shown here: $B = I$.

$$g_\mu(x) = \frac{f(x + \mu u) - f(x)}{\mu} u$$

$$\hat{g}_\mu(x) = \frac{f(x + \mu u) - f(x - \mu u)}{2\mu} u$$

$$g_0(x) = f'(x; u)u$$

$$\mathbb{E}\|g_0(x)\|^2 \leq (n + 4)\|\nabla f(x)\|^2$$




sample $u \sim \mathcal{N}(0, I)$, query finite difference, update in a random direction

- Useful direction only in expectation.
- Second moment grows with n .
- This creates the typical dimension cost.

Accelerated Random Search

Paper Sec. 6 — Owner: Leader



Nesterov-style acceleration with a random finite-difference oracle

$$\mathbb{E}[f(x_k)] - f^* = O\left(\frac{n^2 L_1 R^2}{k^2}\right) + \text{smoothing error}.$$

- Zeroth-order analogue of Nesterov acceleration.
- Keeps the $1/k^2$ shape.
- Does not remove the dimension penalty.

Method	Full gradient	Random gradient-free
Basic smooth	$O(1/k)$	$O(n/k)$
Accelerated smooth	$O(1/k^2)$	$O(n^2/k^2)$

Nonsmooth Convex Optimization


Paper Sec. 4 — Owner: Member A

$$\min_{x \in Q} f(x), \quad f \text{ convex and nonsmooth}$$

$$x_{k+1} = \pi_Q(x_k - h_k g_\mu(x_k))$$

$$\text{For } RS_0: \quad \mathbb{E}[f(\hat{x}_N)] - f^* = O\left(\frac{L_0 R \sqrt{n}}{\sqrt{N}}\right)$$

- Same spirit as projected subgradient descent.
- Random direction replaces a chosen subgradient.



projection returns the random step to the feasible set

Finite Differences and Smoothing Bias

Paper Sec. 4 — Owner: Member A

Method	Oracle	Message
RS_0	directional derivative	clean theory
RS_mu	forward finite difference	true zero-order
subgradient	subgradient	needs subgradient

$$f_\mu(x^*) - f^* \leq \mu L_0 \sqrt{n}$$

$RS_0 : O(n/\varepsilon^2), \quad RS_\mu : O(n^2/\varepsilon^2)$ in the nonsmooth finite-difference setting.


- Smaller μ reduces smoothing bias.
- Too small μ can increase numerical cancellation.
- Theory gives safe bounds, not necessarily best empirical tuning.

Stochastic Zeroth-Order Optimization

Paper Sec. 4 — Owner: Member A

$$f(x) = \mathbb{E}_{\xi}[F(x, \xi)]$$

$$s_{\mu}(x) = \frac{F(x + \mu u, \xi) - F(x, \xi)}{\mu} u$$



Paper summary rate shape: $O\left(\frac{n}{\sqrt{k}}\right)$ in expectation.

- Randomness comes from both u and ξ .
- Reusing ξ in the difference reduces sampling noise.

Smooth Convex Baseline: Random Gradient Descent

Paper Sec. 5 — Owner: Member B

$$\min_x f(x), \quad f \in C^{1,1}$$

$$x_{k+1} = x_k - hg_\mu(x_k), \quad h = \frac{1}{4(n+4)L_1}$$

$$\mathbb{E}[f(\hat{x}_N)] - f^* = O\left(\frac{nL_1R^2}{N}\right) + O(\mu^2 n^2 L_1)$$

- Gradient descent: $O(1/k)$.
- Random gradient-free: still $O(1/k)$, but with a dimension factor.
- This is the theory bridge to the quadratic-chain experiment.


Nonconvex Extension: What Changes?

Paper Sec. 7 — Owner: Member B

$$\min_x f(x), \quad f \text{ may be nonconvex}$$

Goal: make $\|\nabla f(x)\|$ small, not prove global optimality.

- Convex objective gap is no longer the right measure.
- Analyze expected gradient norm or ∇f_μ .
- The target is a stationary point.



Smooth Quadratic Experiment: Setup

Paper Sec. 8.1 — Owner: Member B

$$f_n(x) = \frac{1}{2}x_1^2 + \frac{1}{2} \sum_{i=1}^{n-1} (x_{i+1} - x_i)^2 + \frac{1}{2}x_n^2 - x_1$$

- Smooth quadratic chain from paper Sec. 8.1.
- Known optimizer, so objective gaps are computable.
- Methods: GM, RG_μ , diagnostic RG_0 .
- Dimensions: 32, 64, 128; seeds: 1:10.


Metrics

Relative objective gap, blocks of n random directions, and function-call accounting.

$$\mu = \max\{\mu_{\text{theory}}, 10^{-8}\}$$

Smooth Quadratic Experiment: Results

Paper Sec. 8.1 — Owner: Member B



Method	n	Normalized budget to 10^{-3}	Success
GM	32	3.9	100%
RG_μ	32	—	0%
GM	64	1.0	100%
RG_μ	64	270.0	100%
GM	128	0.2	100%
RG_μ	128	83.2	100%

RG rows count blocks of n random directions; GM rows report

iterations divided by n .

- Full-gradient GM reaches 10^{-3} much earlier.
- RG_μ decreases steadily and nearly matches diagnostic RG_0 .
- Since one RG block uses n random directions, this supports the dimension-cost story; it is not a proof.

Final Takeaways and Limitations

Paper Sec. 8.4 — Owner: Member B

- Gaussian smoothing makes finite-difference random vectors estimate ∇f_μ in expectation.
- Random gradient-free methods mimic first-order methods with extra dimension cost.
- Acceleration is possible: smooth convex rate shape $O(n^2/k^2)$.
- Use these methods when gradients are unavailable, unreliable, or too costly.

Setting	Classical method shape	Random gradient-free shape
Nonsmooth convex	$O(1/\sqrt{k})$	$O(\sqrt{n/k})$ or worse for finite differences
Smooth convex	$O(1/k)$	$O(n/k)$
Accelerated smooth	$O(1/k^2)$	$O(n^2/k^2)$
Nonconvex smooth	Stationarity	Expected gradient-norm control

Limitations

Worse high-dimensional scaling, sensitivity to μ , and slower than accurate gradient methods when gradients are available.

Backup B1: Unified Notation

Symbol	Meaning
n	dimension of the decision variable
$f(x)$	original objective function
$f_\mu(x)$	Gaussian-smoothed objective, $\mathbb{E}[f(x + \mu u)]$
u	Gaussian random direction, usually $u \sim \mathcal{N}(0, I)$
μ	smoothing or finite-difference radius
$g_\mu(x)$	forward finite-difference random oracle
$\hat{g}_\mu(x)$	symmetric finite-difference random oracle
$g_0(x)$	directional-derivative oracle
L_0, L_1	Lipschitz constants for nonsmooth and smooth cases
R	initial distance bound, typically $\ x_0 - x^*\ $
Q, π_Q	feasible set and Euclidean projection

For presentation we use the Euclidean case $B = I$; the paper states the general geometry with Bu in the oracle and $B^{-1}g$ in the update.

Backup B2: Gaussian Gradient Identity

$$f_\mu(x) = \frac{1}{\mu^n \kappa} \int f(y) \exp\left(-\frac{\|y - x\|^2}{2\mu^2}\right) dy.$$

- Change variables with $y = x + \mu u$.
- Differentiate under the integral.
- Use $\mathbb{E}[u] = 0$ to subtract $f(x)$ inside the expectation.

$$\nabla f_\mu(x) = \mathbb{E}_u \left[\frac{f(x + \mu u) - f(x)}{\mu} u \right].$$

Backup B3: Accelerated Method Skeleton

- 1 Maintain two sequences: current point x_k and auxiliary point v_k .
- 2 Form an intermediate point y_k from x_k and v_k .
- 3 Query a random oracle $g_\mu(y_k)$.
- 4 Update both x_{k+1} and v_{k+1} using the oracle.

$$\mathbb{E}[f(x_k)] - f^* = O\left(\frac{n^2 L_1 R^2}{k^2}\right) + \text{smoothing error}.$$

Presentation rule

Show the structure and rate on the main slide; keep the parameter algebra here.

Backup B4: Function Call Accounting

- g_μ usually needs $f(x)$ and $f(x + \mu u)$, unless $f(x)$ is cached.
- $\hat{g}_\mu$ needs two shifted function values.
- Paper tables often report blocks of n random iterations.
- A gradient evaluation is not universally comparable to one function call.

Backup B5: Dimension-Scaling Summary

Method	n	Normalized budget to 10^{-3}	Success
GM	32	3.9	100%
RG_μ	32	—	0%
GM	64	1.0	100%
RG_μ	64	270.0	100%
GM	128	0.2	100%
RG_μ	128	83.2	100%

- The table reports normalized budget to reach relative gap 10^{-3} .
- RG_μ rows count blocks of n random directions; GM rows report iterations divided by n .
- GM uses full gradients, so its cost is not directly comparable to function calls.
- RG_μ is the function-value random method; RG_0 is diagnostic and appears in the detailed CSV.